

Wai Tong Chung

Email: wt.chung94@gmail.com

Personal Web: waitong94.github.io

Experience

Amazon

Member of Technical Staff, AGI Autonomy

San Francisco, CA

Dec 2025 - Present

- Owned continual pre-training, SFT, and RL recipes and infrastructure for agentic large language models (LLMs).
- Achieved 91% on AIME 2026, 68% on SWE-bench Verified and 45% on Terminal-Bench 2.0 with a 100B MoE model.
- Core maintainer of the internal Miles-based trainer within a ~100-person research org.

Together AI

AI Researcher

San Francisco, CA

July 2024 - Dec 2025

- Owned pre- and post-training recipe and stack for inference optimization and agentic language-model applications.
- Developed speculator model training methods that resulted in the fastest B200 inference of DeepSeek-R1 (July 2025).
- Authored four technical blogs on speculative decoding for reasoning-model inference optimization.

Stanford University

Machine Learning Graduate Research Assistant

Stanford, CA

Sept. 2018 - June 2024

- Investigated and developed deep learning methods for efficient high-performance computing software in flow physics.
- Authored 20+ AI for science and computational engineering refereed publications in top ML and engineering venues.
- Wrote accepted NSF, NASA, U.S. DoE, and Google grant proposals (total worth > \$1.5M).

JPMorgan Chase & Co.

Financial Messaging Software Engineer

United Kingdom

Sept. 2017 - Aug. 2018

- Developed, tested, and deployed a Java-based financial messaging application that processed \$6T of daily payments.

Education

Stanford University

Ph.D. in Mechanical Engineering; AI for Science and High-Performance Computing.

Stanford, CA

Sept. 2018 - June 2024

Imperial College London

B.Eng. M.Eng. in Mechanical Engineering, First Class Honours; Best Thesis Prize.

United Kingdom

Sept. 2013 - Aug. 2017

Selected Writing

See [Google Scholar](#) for full list of academic publications.

Technical Blogs

W.T. Chung et al. Boosting DeepSeek-R1's Speed with Customized Speculative Decoding. *Together AI Blog*, 2025. [\[url\]](#)

J. Wang, **W.T. Chung** et al. Together AI Delivers Fastest Inference for the Top Open-Source Models. *Together AI Blog*, 2025. [\[url\]](#), [\[press\]](#)

S. Zhu, F. Bianchi, **W.T. Chung** et al. How Together AI Uses AI Agents to Automate Complex Engineering Tasks: Lessons from Developing Efficient LLM Inference Systems. *Together AI Blog*, 2025. [\[url\]](#)

J. Wang, . . . , **W.T. Chung** et al. AdapTive-LeArning Speculator System (ATLAS): A New Paradigm in LLM Inference via Runtime-Learning Accelerators. *Together AI Blog*, 2025. [\[url\]](#), [\[press\]](#)

Together AI[†]. Together AI Delivers Top Speeds for DeepSeek-R1-0528 Inference on NVIDIA Blackwell. *Together AI Blog*, 2025. ([†]**W.T. Chung** contributed to speculative decoding and benchmarking). [\[url\]](#)

Research Articles

C.-A. Oncescu, Q. Wu, **W.T. Chung** *et al.* Opportunistic Expert Activation: Batch-Aware Expert Routing for Faster Decode Without Retraining. *Int. Conf. Mach. Learn. (ICML)*, 2026. [[.pdf](#)]

M. Ihme[†], **W.T. Chung**[†]. Artificial Intelligence as a Catalyst for Combustion Science and Engineering[‡]. Accepted in *Proc. Combust. Inst.* 40, 2024. ([†]*Equal contribution*. [‡]Presented as a plenary lecture at the 40th International Symposium on Combustion, Milan, 2024) [[.pdf](#)]

W.T. Chung *et al.* Turbulence in Focus: Benchmarking Scaling Behavior of 3D Volumetric Super-Resolution with BLASTNet 2.0 Data. *Adv. Neural Inf. Process. Syst. (NeurIPS)* 36, 2023. [[.pdf](#), [press](#)]

Selected Professional Activities

Contributor to Miles [[github](#)]

Reviewer for *ICLR* 2025–2026; *NeurIPS* 2024–2026; *AISTATS* 2025; and other AI-for-science venues.

Lead Organizer for *Stanford HAI Climate-Centered AI Seminar Series*, 2023. [[press](#)]

Skills

Programming and Engineering

Training/Inference: PyTorch, Miles, Megatron-LM, SGLang, vLLM.

Systems/Infrastructure: Kubernetes, Slurm, Ray, MPI.

Programming: Python, C++, Java

Languages

Proficient: English, Malay.

Familiar: Mandarin, Cantonese.

Honors and Awards

Stanford CS323: The AI Awakening Best Final Project Prize (Top 4 of 87 Students)	2023
Stanford Human-Centered AI Affinity Group Award [info , press]	2023
Stanford Human-Centered AI Graduate Fellowship [info , press]	2022-2023
Stanford School of Engineering Graduate Fellowship	2018-2019
Imperial College Mechanical Engineering Most Outstanding Thesis Prize (Top 1 of 138 Students)	2017
Imperial College Mechanical Engineering Dean's List (Top 10% of 138 Students)	2017