

# Wai Tong Chung

Email: [wt.chung94@gmail.com](mailto:wt.chung94@gmail.com)

Personal Web: [waitong94.github.io](https://waitong94.github.io)

## Experience

### Amazon

Member of Technical Staff, AGI Autonomy

San Francisco, CA

Dec 2025 - Present

- Owned continual pre-training, SFT, and RL recipes and infrastructure for agentic large language models (LLMs).
- Achieved 92% on AIME 2026, 68% on SWE-bench Verified and 45% on Terminal-Bench 2.1 with a 106B MoE model.
- Core maintainer of the internal Miles-based trainer within a ~100-person research org.

### Together AI

AI Researcher

San Francisco, CA

July 2024 - Dec 2025

- Owned pre- and post-training recipe and stack for inference optimization and agentic language-model applications.
- Developed speculator model training methods that resulted in the fastest B200 inference of DeepSeek-R1 (July 2025).
- Authored four technical blogs on speculative decoding for reasoning-model inference optimization.

### Stanford University

Machine Learning Graduate Research Assistant

Stanford, CA

Sept. 2018 - June 2024

- Investigated and developed deep learning methods for efficient high-performance computing software in flow physics.
- Authored 20+ AI for science and computational engineering refereed publications in top ML and engineering venues.
- Wrote accepted NSF, NASA, U.S. DoE, and Google grant proposals (total worth > \$1.5M).

### JPMorgan Chase & Co.

Financial Messaging Software Engineer

United Kingdom

Sept. 2017 - Aug. 2018

- Developed, tested, and deployed a Java-based financial messaging application that processed \$6T of daily payments.

## Education

### Stanford University

Ph.D. in Mechanical Engineering; AI for Science and High-Performance Computing.

Stanford, CA

Sept. 2018 - June 2024

### Imperial College London

B.Eng. M.Eng. in Mechanical Engineering, First Class Honours; Best Thesis Prize.

United Kingdom

Sept. 2013 - Aug. 2017

## Selected Writing

See [Google Scholar](#) for full list of academic publications.

### Technical Blogs

**W.T. Chung** et al. Boosting DeepSeek-R1's Speed with Customized Speculative Decoding. *Together AI Blog*, 2025. [\[url\]](#)

J. Wang, **W.T. Chung** et al. Together AI Delivers Fastest Inference for the Top Open-Source Models. *Together AI Blog*, 2025. [\[url\]](#), [\[press\]](#)

S. Zhu, F. Bianchi, **W.T. Chung** et al. How Together AI Uses AI Agents to Automate Complex Engineering Tasks: Lessons from Developing Efficient LLM Inference Systems. *Together AI Blog*, 2025. [\[url\]](#)

J. Wang, . . . , **W.T. Chung** et al. AdapTive-LeArning Speculator System (ATLAS): A New Paradigm in LLM Inference via Runtime-Learning Accelerators. *Together AI Blog*, 2025. [\[url\]](#), [\[press\]](#)

Together AI<sup>†</sup>. Together AI Delivers Top Speeds for DeepSeek-R1-0528 Inference on NVIDIA Blackwell. *Together AI Blog*, 2025. (<sup>†</sup>**W.T. Chung** contributed to speculative decoding and benchmarking). [\[url\]](#)

## Research Articles

C.-A. Oncescu, Q. Wu, **W.T. Chung** *et al.* Opportunistic Expert Activation: Batch-Aware Expert Routing for Faster Decode Without Retraining. *Int. Conf. Mach. Learn. (ICML)*, 2026. [[.pdf](#)]

M. Ihme<sup>†</sup>, **W.T. Chung**<sup>†</sup>. Artificial Intelligence as a Catalyst for Combustion Science and Engineering<sup>‡</sup>. Accepted in *Proc. Combust. Inst.* 40, 2024. (<sup>†</sup>*Equal contribution*. <sup>‡</sup>Presented as a plenary lecture at the 40<sup>th</sup> International Symposium on Combustion, Milan, 2024) [[.pdf](#)]

**W.T. Chung** *et al.* Turbulence in Focus: Benchmarking Scaling Behavior of 3D Volumetric Super-Resolution with BLASTNet 2.0 Data. *Adv. Neural Inf. Process. Syst. (NeurIPS)* 36, 2023. [[.pdf](#), [press](#)]

## Selected Professional Activities

**Contributor** to Miles [[github](#)]

**Reviewer** for *ICLR* 2025–2026; *NeurIPS* 2024–2026; *AISTATS* 2025; and other AI-for-science venues.

**Lead Organizer** for *Stanford HAI Climate-Centered AI Seminar Series*, 2023. [[press](#)]

## Skills

### Programming and Engineering

Training/Inference: PyTorch, Miles, Megatron-LM, SGLang, vLLM.

Systems/Infrastructure: Kubernetes, Slurm, Ray, MPI.

Programming: Python, C++, Java

### Languages

Proficient: English, Malay.

Familiar: Mandarin, Cantonese.

## Honors and Awards

|                                                                                                  |           |
|--------------------------------------------------------------------------------------------------|-----------|
| Stanford <a href="#">CS323: The AI Awakening</a> Best Final Project Prize (Top 4 of 87 Students) | 2023      |
| Stanford Human-Centered AI Affinity Group Award [ <a href="#">info</a> , <a href="#">press</a> ] | 2023      |
| Stanford Human-Centered AI Graduate Fellowship [ <a href="#">info</a> , <a href="#">press</a> ]  | 2022-2023 |
| Stanford School of Engineering Graduate Fellowship                                               | 2018-2019 |
| Imperial College Mechanical Engineering Most Outstanding Thesis Prize (Top 1 of 138 Students)    | 2017      |
| Imperial College Mechanical Engineering Dean's List (Top 10% of 138 Students)                    | 2017      |